

Computer Networks - II

Programming assignment - I

Saurabh (CS14BTECH11031)

Part - I

Client-side

The client side program invokes two system calls, the first, “**sendto()**”, where the client sends a packet to the server, and the second, “**recvfrom()**”, where, the same packet is echoed by the server back to the client. The following is the structure of the packet :

```
struct timestamp {  
    struct timeval current_time;  
    long elapsed;  
    char buffer[1024];  
};
```

The first field is the timestamp, the second stores the time elapsed for a one way trip, and the last argument is the character buffer, which stores the time at which the packet was deployed. The socket is created by the client as follows :

```
if((socket_id = socket(AF_INET, SOCK_DGRAM, 0)) == -1){  
    std::cerr << "Socket couldn't be created.\n";  
    return 1;  
}  
server_addr.sin_family = AF_INET;  
server_addr.sin_port = htons(5000);  
server_addr.sin_addr.s_addr = inet_addr("127.0.0.1");  
addr_size = sizeof(struct sockaddr);
```

We can see that the port is set to **5000** and the IP is localhost (which can be changed in case of remote server), and the server protocol is set to UDP using **SOCK_DGRAM**.

Now, coming to the system calls involved in sending and receiving the packets, the `sendto()` system call is used to transmit a message to another socket.

The `send()` call may be used only when the socket is in a connected state (so that the intended recipient is known). On success, this call return the number of bytes sent. On error, -1 is returned, and `errno` is set appropriately.

The **`recvfrom()`** call is used to receive messages from a socket. It may be used to receive data on both connectionless and connection-oriented sockets. If no messages are available at the socket, the receive calls wait for a message to arrive, unless the socket is nonblocking. **`recvfrom()`** places the received message into the buffer `buf`. The caller needs to specify the size of the buffer in `len` argument. On success, returns number of bytes received, on error, -1 is returned and `errno` is set appropriately.

The following is the structure of the invocation of the above two calls in the client program :

```
for(int i = 0; i < num_messages; i++){
    struct timestamp rtt_struct, new_rtt_struct;
    gettimeofday(&rtt_struct.current_time, 0);
    rtt_struct.elapsed = 0;
    currentTime(rtt_struct.buffer);
    int bytes_sent = sendto(socket_id, &rtt_struct, sizeof(rtt_struct), 0, (struct sockaddr *) &server_addr, addr_size);
    if(setsockopt(socket_id, SOL_SOCKET, SO_SNDTIMEO, (char *) &tv, sizeof(struct timeval)) < 0){
        perror("Error in setting socket options\n");
    }
    int bytes_rcvd = recvfrom(socket_id, &new_rtt_struct, sizeof(new_rtt_struct), 0, NULL, NULL);
```

The above program runs a loop and sends/receives the packets, and timeout is set appropriately in the `setsockopt()` method to indicate that the packet has been lost.

Server-side

The server does the opposite sequence as done by the client. The server first waits for a message to be sent from the client, and does a 'blocking' **`recvfrom()`** unless a message from the client is sent to the server, thus, confirming a message/packet sending entity exists, and then upon successful receival, the server "echoes" back the packet to the client and this continues until the client keeps on sending messages/packets. The prerequisites like the setting up of socket, getting a socket-id, setting fields like **`sin_family`**, the IP and port is similar and the fields used are the same as in the client. The following is the additional job needed to be done in the server :

```

addr_size = sizeof(struct sockaddr);
if(bind(socket_id, (struct sockaddr *)&server_addr, addr_size) == -1){
    std::cerr << "Error in binding the socket as socket already in use.\n";
    return 1;
}

```

The server needs to do this because when a socket is created with the socket system call as described in the client section, it exists in a name space (address family) but has no address assigned to it. **bind()** assigns the address specified by **server_addr** to the socket referred to by the file descriptor **socket_id**. **addrlen** (**sizeof server_addr**) specifies the size, in bytes, of the address structure pointed to by **server_addr**. More informally, this operation can be termed as “assigning a name to a socket”.

After the **bind()** operation is complete, a for loop runs which keeps receiving and echoing back the messages to the client as shown in the code :

```

for(int i = 0; i < num_messages; i++){
    struct timestamp rtt_struct;
    int bytes_rcvd = recvfrom(socket_id, &rtt_struct, sizeof(rtt_struct), 0, (struct sockaddr *) &client_addr, &addr_size);
    get_time = rtt_struct.current_time;
    cout << "Received at server : ";
    cout << rtt_struct.buffer << endl;
    gettimeofday(&rtt_struct.current_time, 0);
    rtt_struct.elapsed = (rtt_struct.current_time.tv_sec - get_time.tv_sec) * 1000000
        + rtt_struct.current_time.tv_usec - get_time.tv_usec;
    int bytes_sent = sendto(socket_id, &rtt_struct, sizeof(rtt_struct), 0, (struct sockaddr *) &client_addr, addr_size);
}

```

Summarizing what happens at the client-side and the server-side, we have :

At client-side :

0. **gethostbyname()**
1. **socket()**
2. **memset(server_addr.sin_zero, '\0', sizeof(server_addr.sin_zero))**
4. **sendto()**
5. **recvfrom()**

And at server-side :

0. **Variable initialization**

- 1.socket()
- 2.bind()
- 3.recvfrom()
- 4.sendto()

Part - II

Client-side

The only change implemented in this code, is that a while loop is run until the interval of sleep method is positive, we keep sending packets, and inside the while loop, we further run a for loop, with maximum limit set to a variable which increases by 1 in every while loop iteration. At the end of a while loop, the sleep interval decreases by 1 second. So basically, in the first iteration, 1 packet is sent, in the second iteration, 2 packets are sent, 3 in the 3rd, and so until the interval between two consecutive batch of packets to be sent is positive. This is implemented as following :

```
while(interval > 0){
    for(int i = 0;i < no_of_packets;i++){
        currentTime(buffer_rcvd);
        gettimeofday(&t0, 0);
        bytes_sent = sendto(socket_id, buffer_rcvd, sizeof(buffer_rcvd), 0, (struct sockaddr *) &server_addr, addr_size);
        buffer_rcvd[bytes_sent] = '\0';
        bytes_rcvd = recvfrom(socket_id, buffer_rcvd, sizeof(buffer_rcvd), 0, (struct sockaddr *) &server_addr, &addr_size);
        gettimeofday(&t1, 0);
        cout << "Received message on trip " << trip << " : " << buffer_rcvd << endl;
        total_size += bytes_sent;
        total_time += ((t1.tv_sec - t0.tv_sec) * 1000000
            + t1.tv_usec - t0.tv_usec);
        total_time /= 1000000.0;
        throughput += (total_size)/total_time;
        average_delay = total_time/count++;
        for(int j = 0;j < interval;j++){
            fprintf(fp, "%lf\n", throughput);
            fprintf(fp_, "%lf\n", average_delay);
        }
    }
    no_of_packets++, trip++;
    sleep(interval--);
}
```

The appropriate **throughput** (bytes/second) and **average_delay** (seconds) calculation is done at the end of loop.

Server-side

The server-side code in 2nd part is mostly same as that in the 1st part. Same `sendto()` and `recvfrom()` system calls are invoked to echo all the messages sent by the client back to it. The following is the main code doing this :

```
if(bind(socket_id, (struct sockaddr *)&server_addr, addr_size) == -1){
    std::cerr << "Error in binding the socket as socket already in use.\n";
    return 1;
}
while(1){
    bytes_rcvd = recvfrom(socket_id, buffer_rcvd, sizeof(buffer_rcvd), 0, (struct sockaddr *) &client_addr, &addr_size);
    buffer_rcvd[bytes_rcvd] = 0;
    cout << "Message received at server : " << buffer_rcvd << endl;
    bytes_sent = sendto(socket_id, buffer_rcvd, sizeof(buffer_rcvd), 0, (struct sockaddr *) &client_addr, addr_size);
}
```

The following is the screenshot of the packet capture using wireshark, where my IP is **172.16.2.23** and the remote host's IP is **172.16.2.41**. The only configuration needed to run a client-server echo program is set the IP address to be the IP address of the host computer in the client-side program :

```
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(5000);
server_addr.sin_addr.s_addr = inet_addr("172.16.2.41");
addr_size = sizeof(struct sockaddr);
```

The following is the screenshot showing my machine's IP address, which is the IP address of the client :

```
wlp6s0    Link encap:Ethernet  HWaddr 4c:bb:58:13:f2:c8
          inet addr:172.16.2.23  Bcast:172.16.7.255  Mask:255.255.248.0
```

The following screenshot clearly shows the echo between the client and the server, as the packets travel back-and-forth :

78	4.389011272	172.16.2.23	172.16.2.41	UDP	1090	56010 → 5000	Len=1048
79	4.389095776	172.16.2.41	172.16.2.23	UDP	1090	5000 → 56010	Len=1048
112	6.402650544	172.16.2.23	172.16.2.41	UDP	1090	56010 → 5000	Len=1048
113	6.402790881	172.16.2.41	172.16.2.23	UDP	1090	5000 → 56010	Len=1048
145	8.418413976	172.16.2.23	172.16.2.41	UDP	1090	56010 → 5000	Len=1048
146	8.418523169	172.16.2.41	172.16.2.23	UDP	1090	5000 → 56010	Len=1048
170	10.434118548	172.16.2.23	172.16.2.41	UDP	1090	56010 → 5000	Len=1048
171	10.434274367	172.16.2.41	172.16.2.23	UDP	1090	5000 → 56010	Len=1048
191	12.450125165	172.16.2.23	172.16.2.41	UDP	1090	56010 → 5000	Len=1048
192	12.450255408	172.16.2.41	172.16.2.23	UDP	1090	5000 → 56010	Len=1048
207	14.463931856	172.16.2.23	172.16.2.41	UDP	1090	56010 → 5000	Len=1048
208	14.464081709	172.16.2.41	172.16.2.23	UDP	1090	5000 → 56010	Len=1048
217	16.479060065	172.16.2.23	172.16.2.41	UDP	1090	56010 → 5000	Len=1048
218	16.479176827	172.16.2.41	172.16.2.23	UDP	1090	5000 → 56010	Len=1048

The graph for **throughput** (bytes/sec) and **average_delay** (sec) is saved separately in the assignment directory along with the screenshots of the working of the code.